



Distributed Systems

CSCI 5673

1/8/2026

Mark Zhao



College of Engineering & Applied Science

UNIVERSITY OF COLORADO BOULDER

What is a distributed system?

Collection of devices

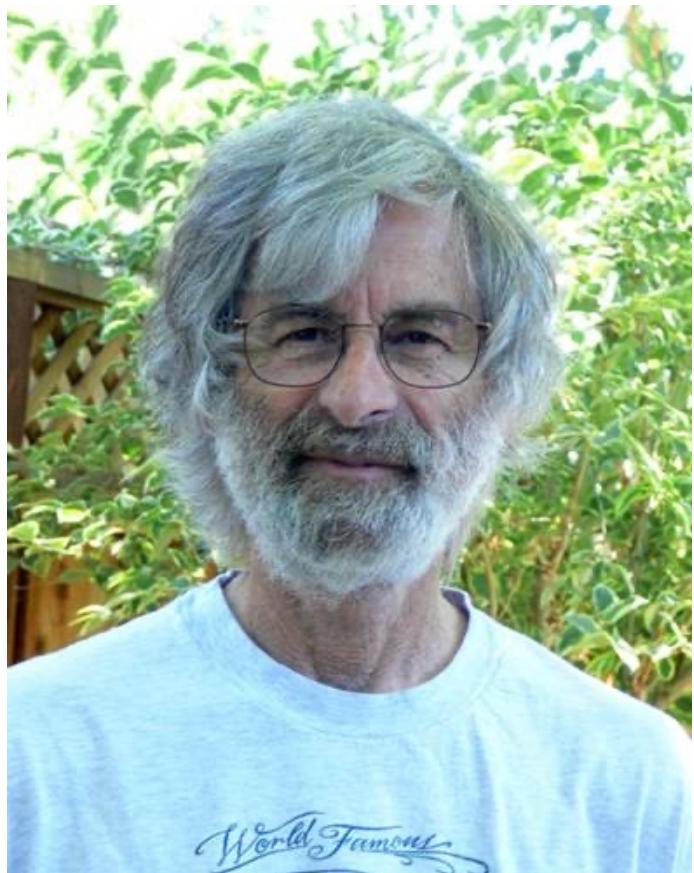
- Collaborating on a task

Internet \supset Networked

- Autonomous

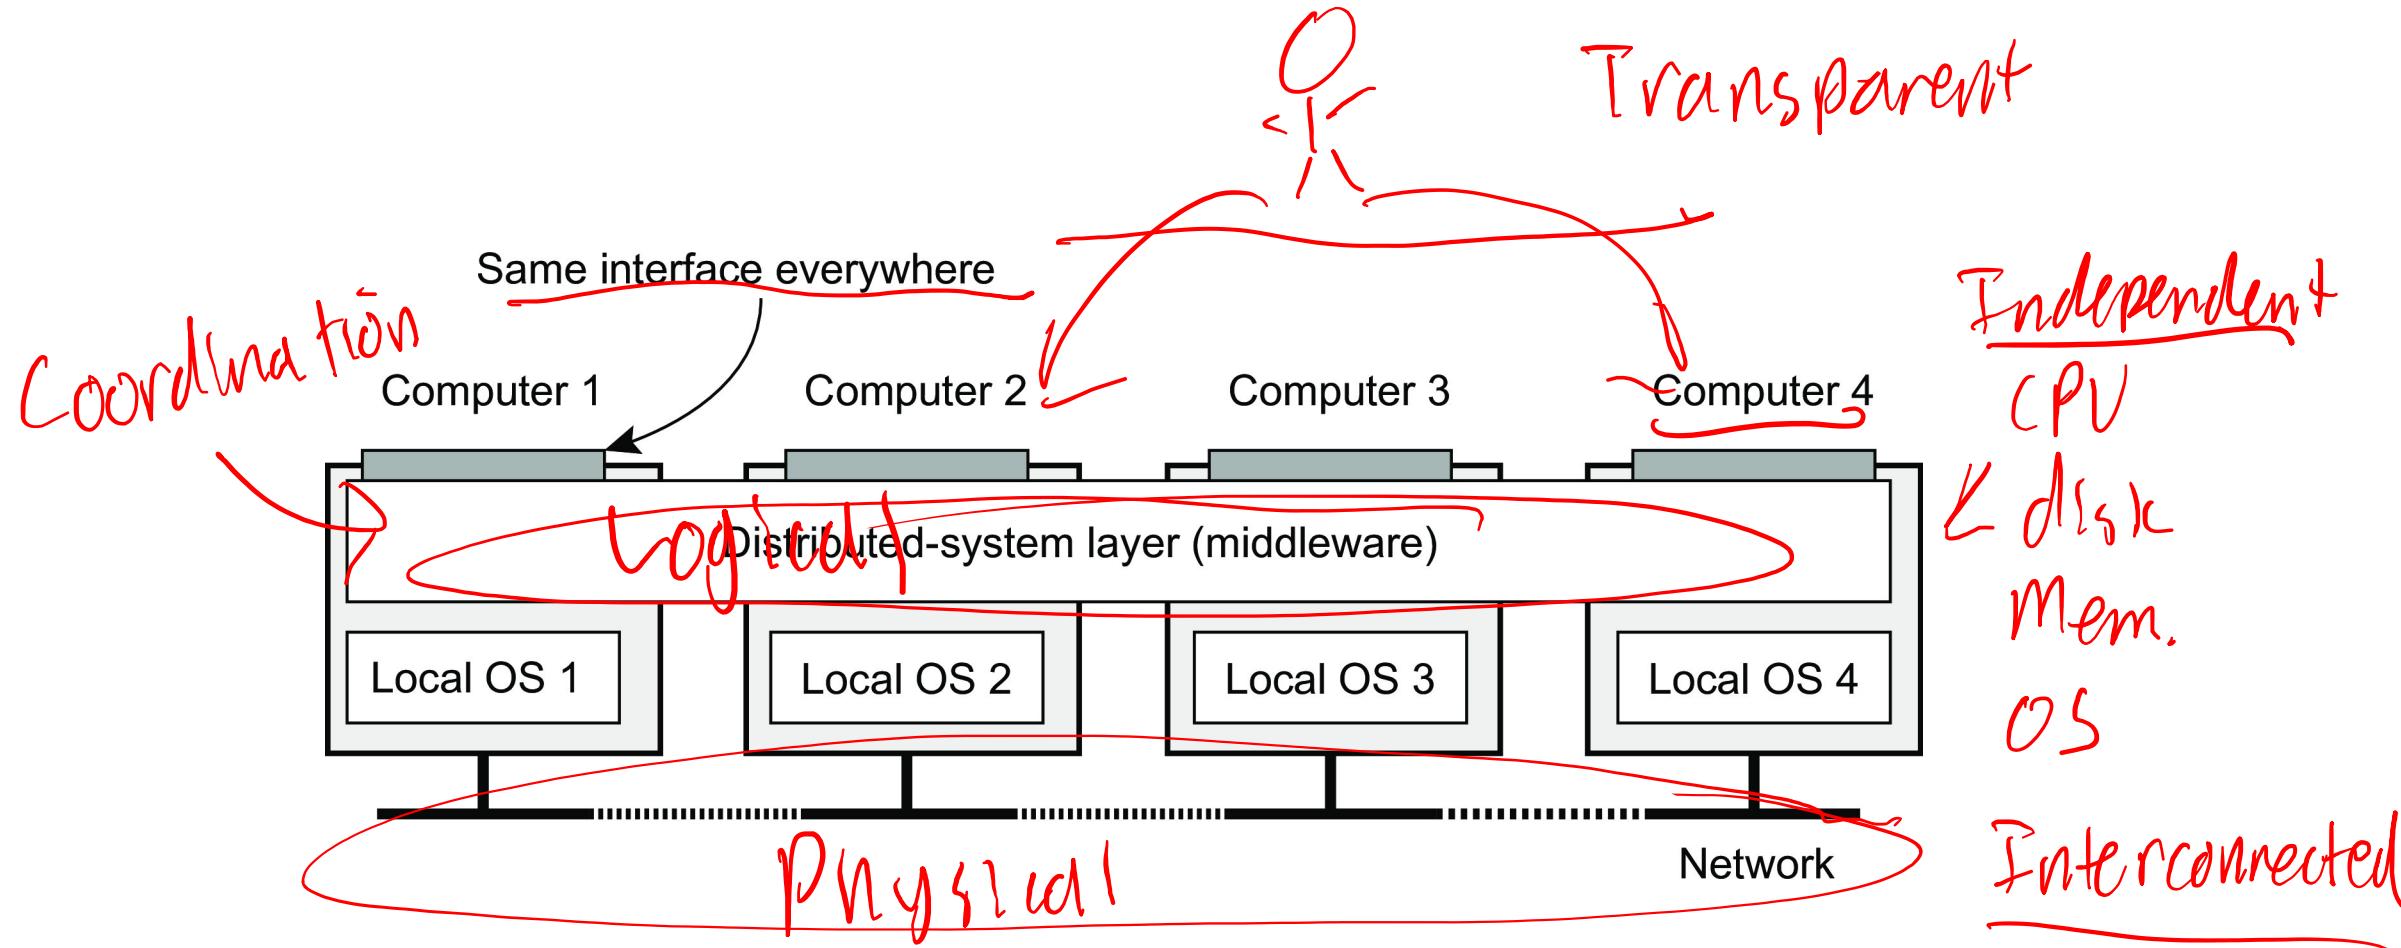
(For the most part)

What is a distributed system?



“... a system in which the failure of a computer you didn’t even know existed can render your own computer unusable.” – Leslie Lamport

What is a distributed system?



What is a distributed system?

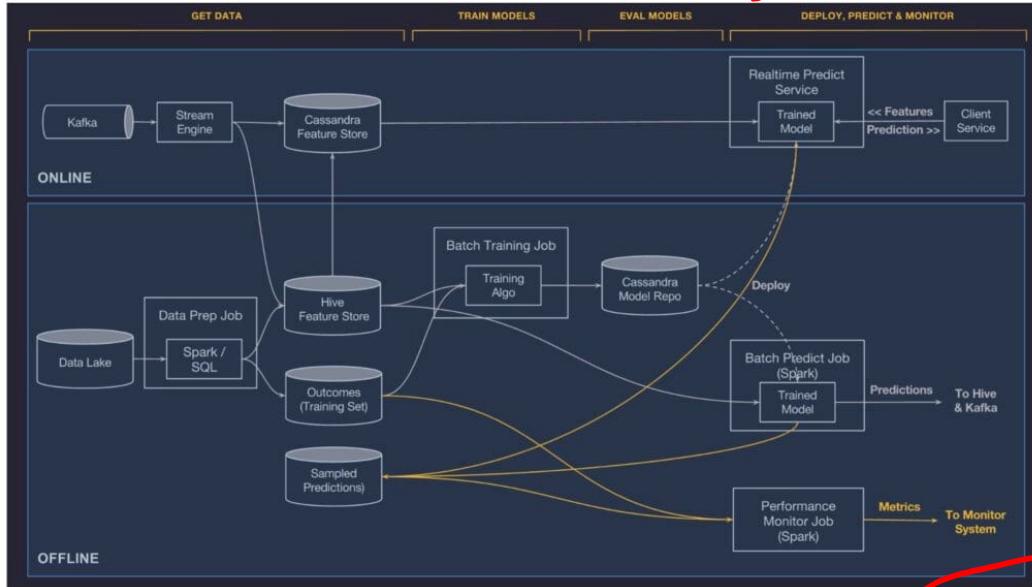
A collection of **independent** computing components, **interconnected** via a network, that are capable of **collaborating** on a task.

- **Independent**
 - Own set of compute resources
- **Interconnected**
 - Communicate over a network to exchange information
- **Collaborative:**
 - No centralized control
 - **Hide** the fact that processes/resources are physically distributed, possibly globally

What is a distributed system?

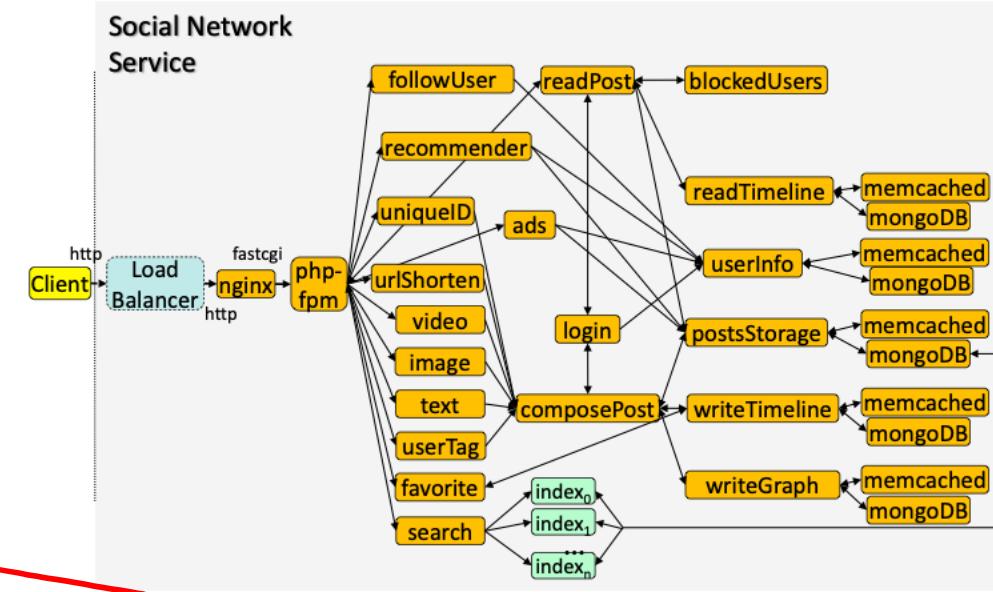
Databases
→ Transactions

Uber's ML System



UIC
OIS
Facebook
US

Micro service



lower Latency
→ CDNs

<https://www.uber.com/blog/michelangelo-machine-learning-platform/>

<https://www.csl.cornell.edu/~delimitrou/papers/2019.asplos.microservices.pdf>

Why do we want to distribute systems?

- Low Latency (CDNs)
- Parallelism / Scalability
 - $10x$ resources \rightarrow $10x$ More throughput / tasks
- Reliability / Availability
 - \Rightarrow If one node fails, we can still make progress
- Features / Modularity / Extensibility

Why do we want to distribute systems?

- Scalability
 - Maintain performance as users/data/... grows
 - How? Increase compute/memory/storage/network/... capacity via parallelism
- Incremental Growth
 - Gradually add/remove components as needed
- Fault tolerance
 - Keep going even if one/few nodes fail
- Distribution
 - Match physical placement of components to application needs

CDNs
Mobile devices
P2P
Blockchain

Why do we want to distribute systems?

- Security
 - Isolate sensitive components from others
 - Modularity, principle of least privilege
 - Mutually distrusting parties *← blockchains*
- Heterogeneity
 - Adapt to different HW, devices, etc.
- Mobility
 - Mobile and “Ubiquitous” computing

Big goal: Hide the complexity of distribution from applications

Why do we not want to distribute systems?

- **Complexity**

- Comm.* • System architecture much more involved
 - Administrative complexity / *cost*

- *Comm* **Performance bottlenecks**

- "Critical path" involves many more steps

- **Inconsistencies**

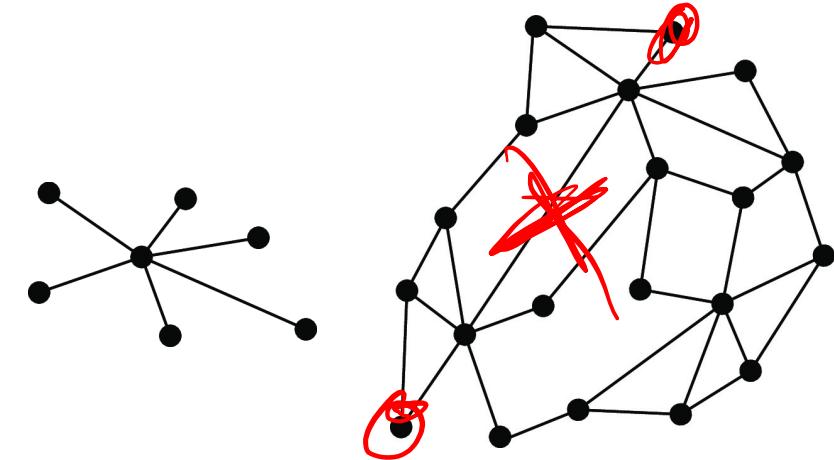
- What if two nodes disagree?

- **Partial failures**

- What if systems/networks fail?

- **Security**

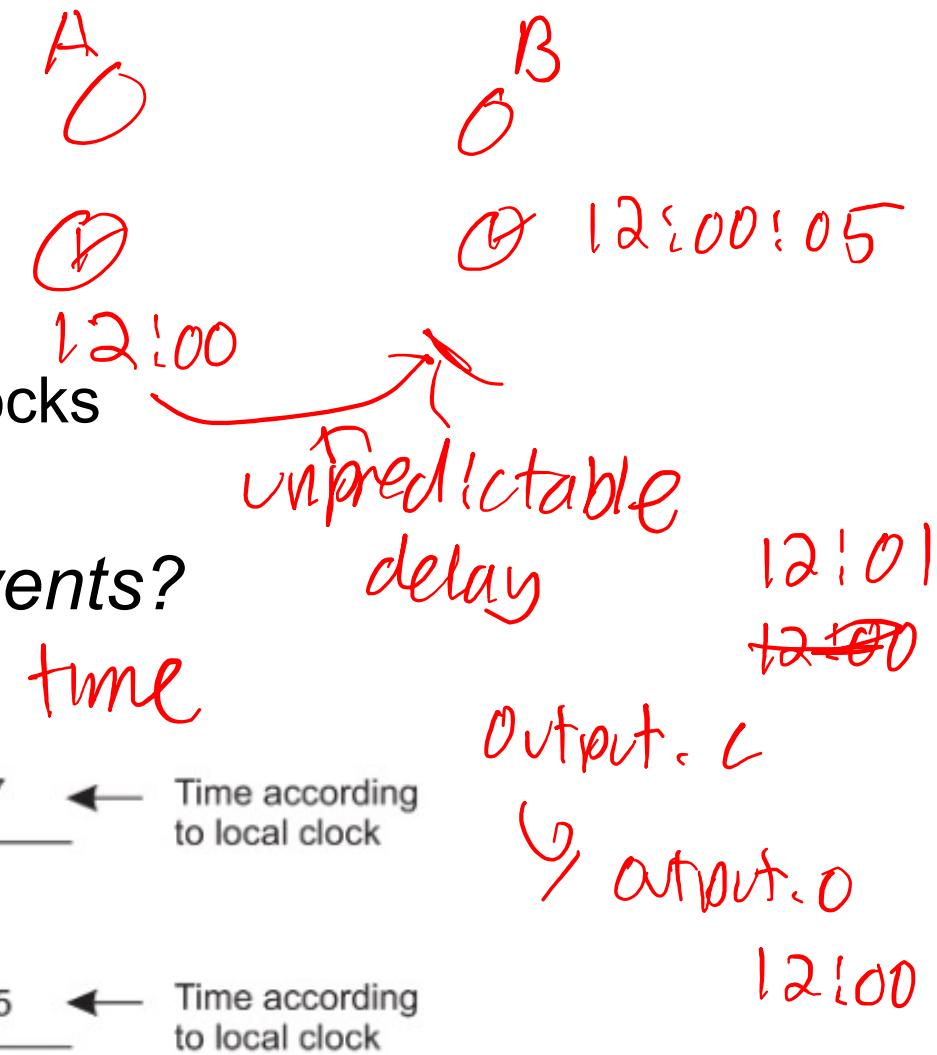
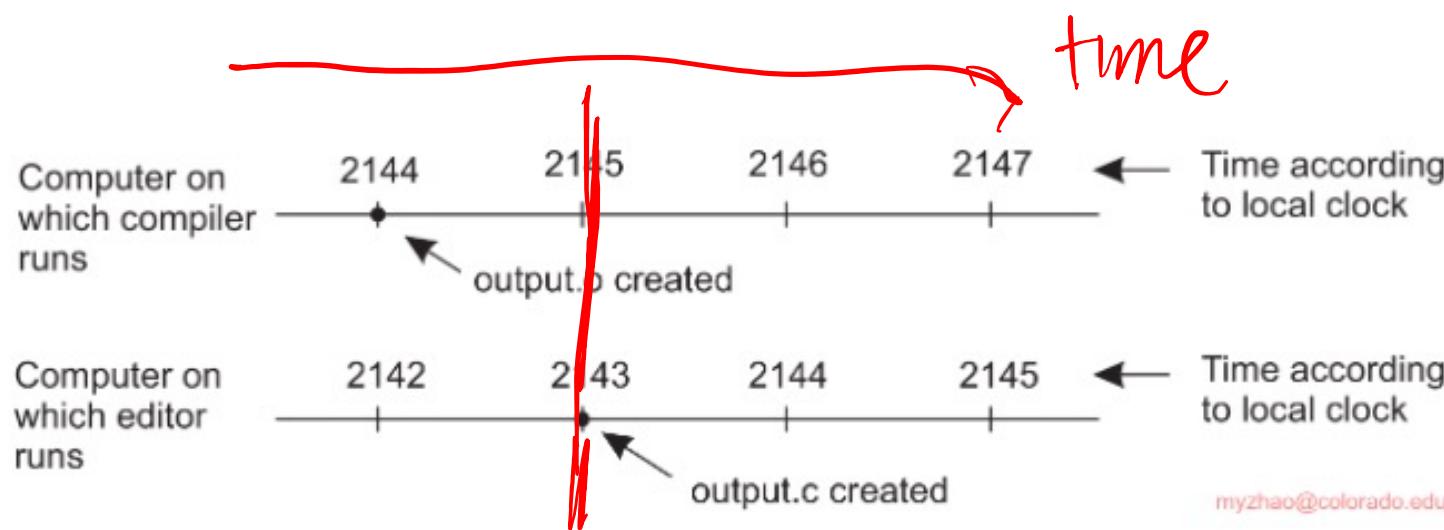
- How do we ensure components are trusted?



(Some) Limitations of distributed systems

Absence of a common global clock

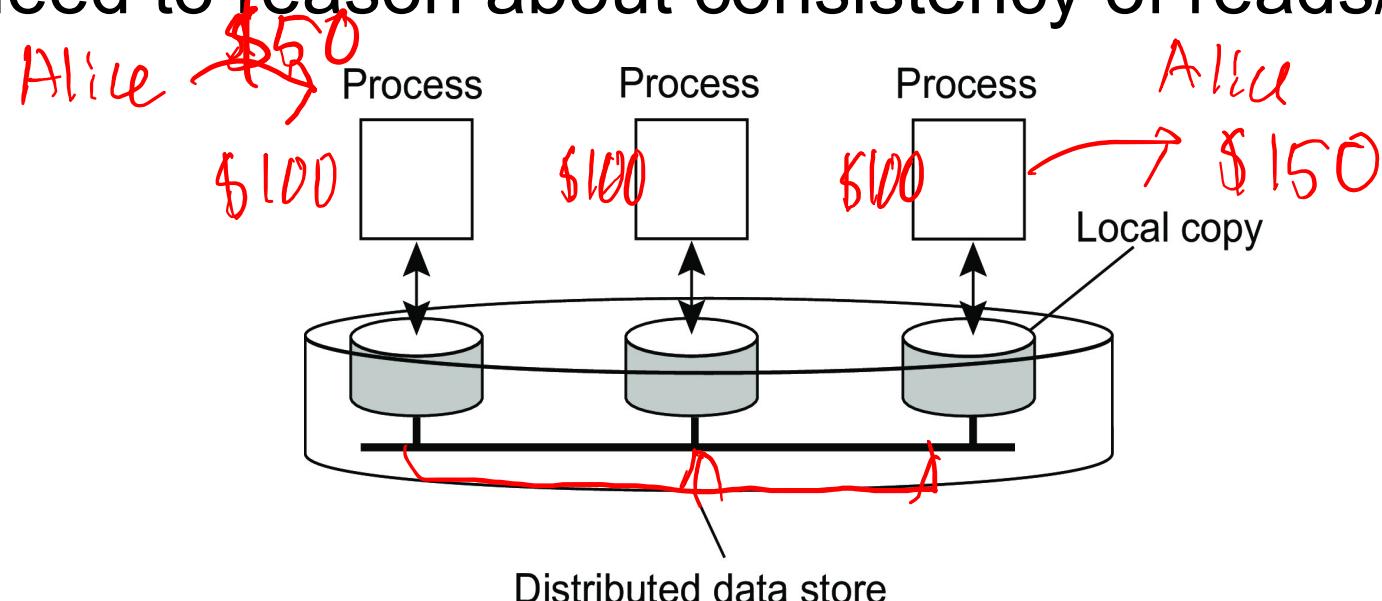
- No notion of a shared “global” time
 - Each component has its own clock
 - Impossible to perfectly synchronize all local clocks
- Challenge: What's the temporal order of events?



(Some) Limitations of distributed systems

Absence of shared memory

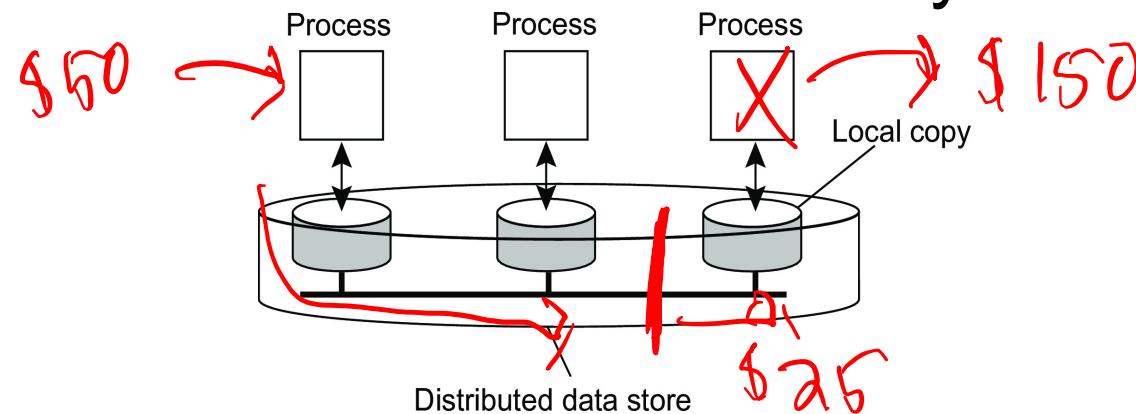
- Notion of a common shared memory (RAM, disk, ...) does not exist
- Each component may have its own memory
- *Challenge:* Need to reason about consistency of reads/writes across replicas



(Some) Limitations of distributed systems

Absence of any guarantees in reliability

- Real-world servers/networks can “break” at any point
 - No guarantees that processes are always alive
 - No guarantees that network messages ever arrive
 - ...and if they do, no guarantees on *when* they arrive
 - No guarantees that messages arrive intact
 - What if a malicious actor or just pure chance causes a message to say the “wrong” thing
- *Challenge:* How do we still let the overall system make progress?



Why study distributed systems?

- These challenges come up when building pretty much any modern application
 - Web apps, AI systems, mobile apps, ...
- Understanding how to reason about and solve these issues will help across a broad set of challenges
 - System design in industry – what tools / techniques to use to deploy your product
 - Active research area – lots of unsolved problems!
- It's interesting!
 - Hard problems with powerful solutions

Teaching staff

- Mark Zhao
 - Assistant professor, computer science
 - <https://basil.colorado.edu>
- Rezwan Rahman
 - PhD candidate, systems
- Doncey Albin
 - PhD candidate, robotics
- Charith Purushotham
 - MS CS

Course structure

- Course website is the main source of truth for schedule/logistics/etc.
 - <https://basil.colorado.edu/courses/sp26/csci5673/about/>
- Canvas
 - Course announcements, assignment releases, lecture slides
- Gradescope
 - Assignment submission
- Piazza
 - Course Q&A and discussion forum

Course structure

- Homeworks and programming assignments (40%)
 - ~4 homework assignments based on concepts from lectures
 - ~4 lab assignments that get you to build a large-scale distributed system from scratch
- Exams (40%)
 - Midterm in class on 3/12
 - Final exam during finals week
- Project (20%)
 - Open-ended project building and deploying a distributed system
 - Intermediate project updates to build up to a full “paper”
 - In groups of up to 3, start forming project groups!

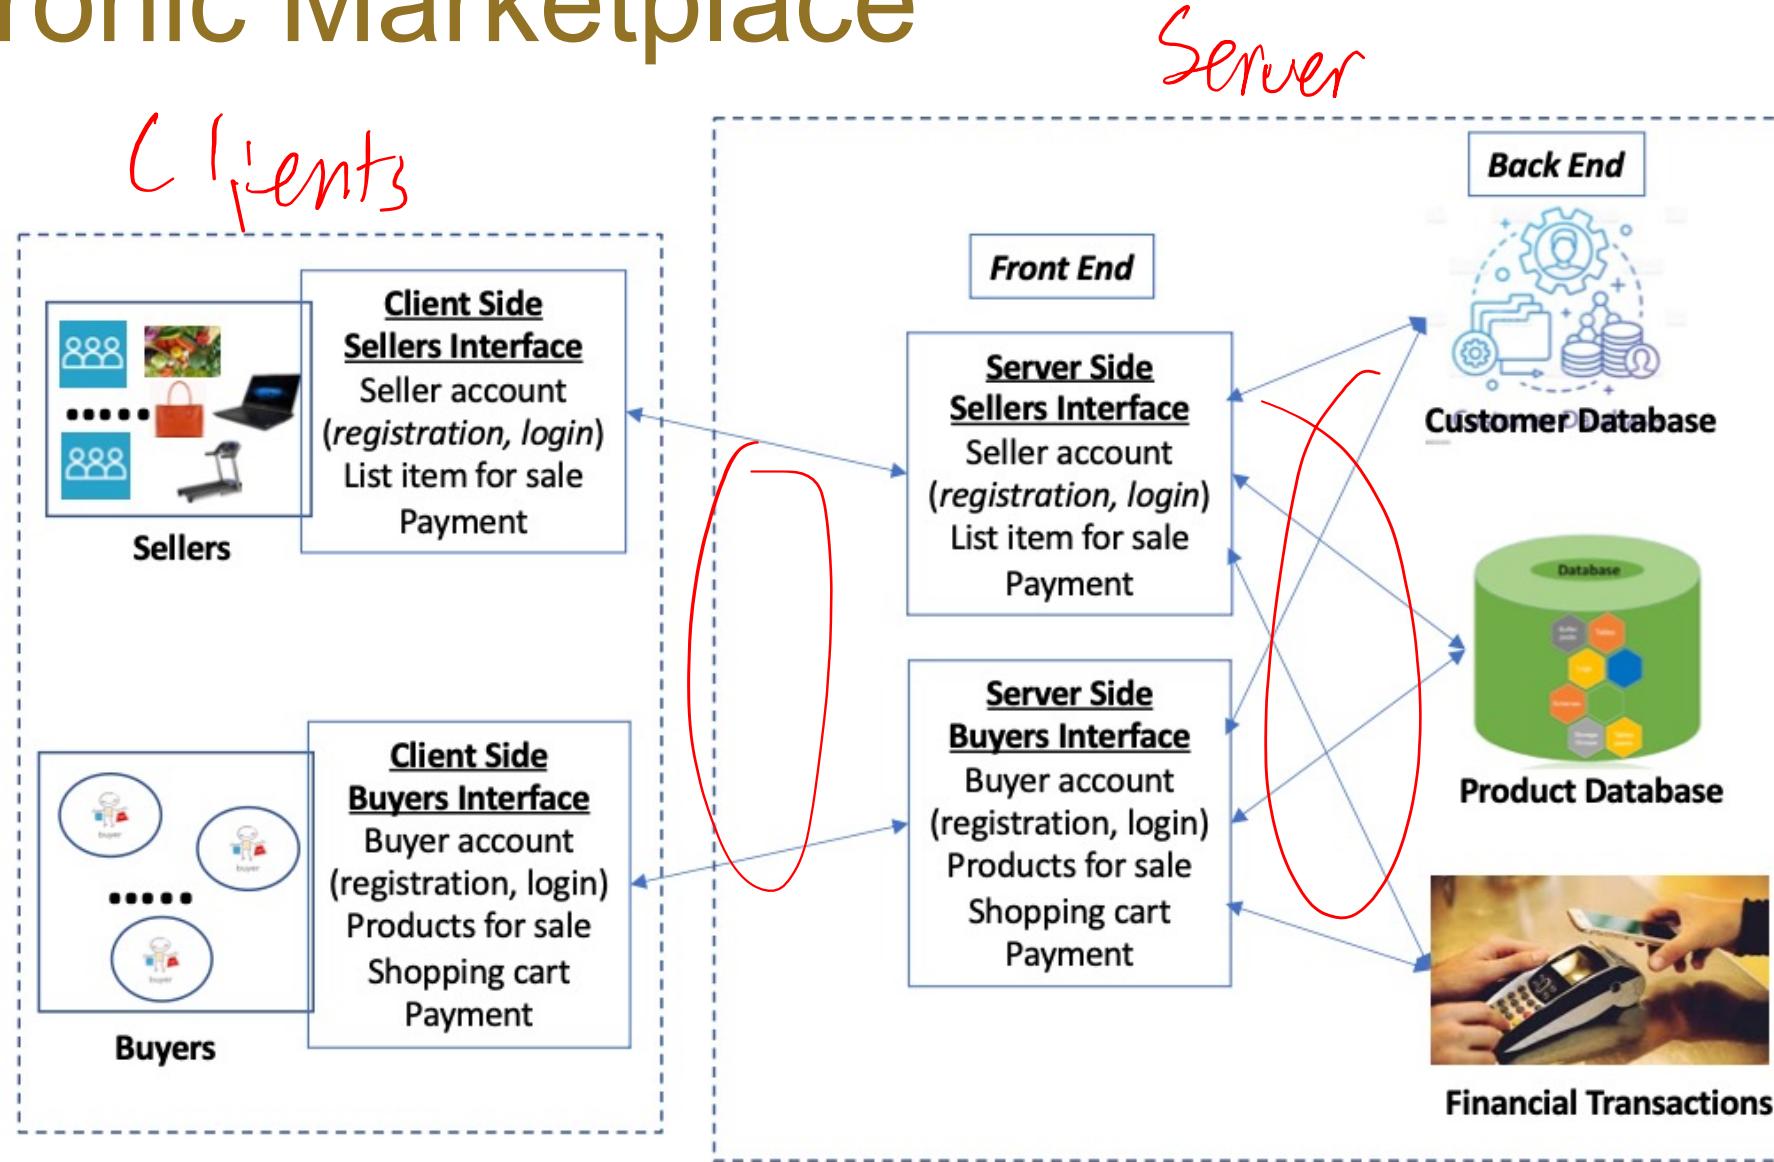
Prereqs and Materials

- Course in Operating Systems (CSCI 3753)
 - Processes, threads, synchronization, memory management, file systems, virtualization, ...
- Network Systems (CSCI 4273/5273)
 - TCP/UDP, HTTP, DNS, Socket API and programming, wide area networks, ...
- Readings published on course schedule
 - Distributed Systems by van Steen and Tanenbaum (available online)
 - Designing Data-Intensive Applications by Kleppmann (available online)
 - Selected papers from literature (mostly around large-scale, distributed systems deployed at Google/Amazon/Meta/...)

Programming assignments / course outline

- Your job is to build up a distributed marketplace system
 - E.g., Amazon, eBay, etc.
- Sellers can list items for sale
- Buyers can view and purchase items for sale

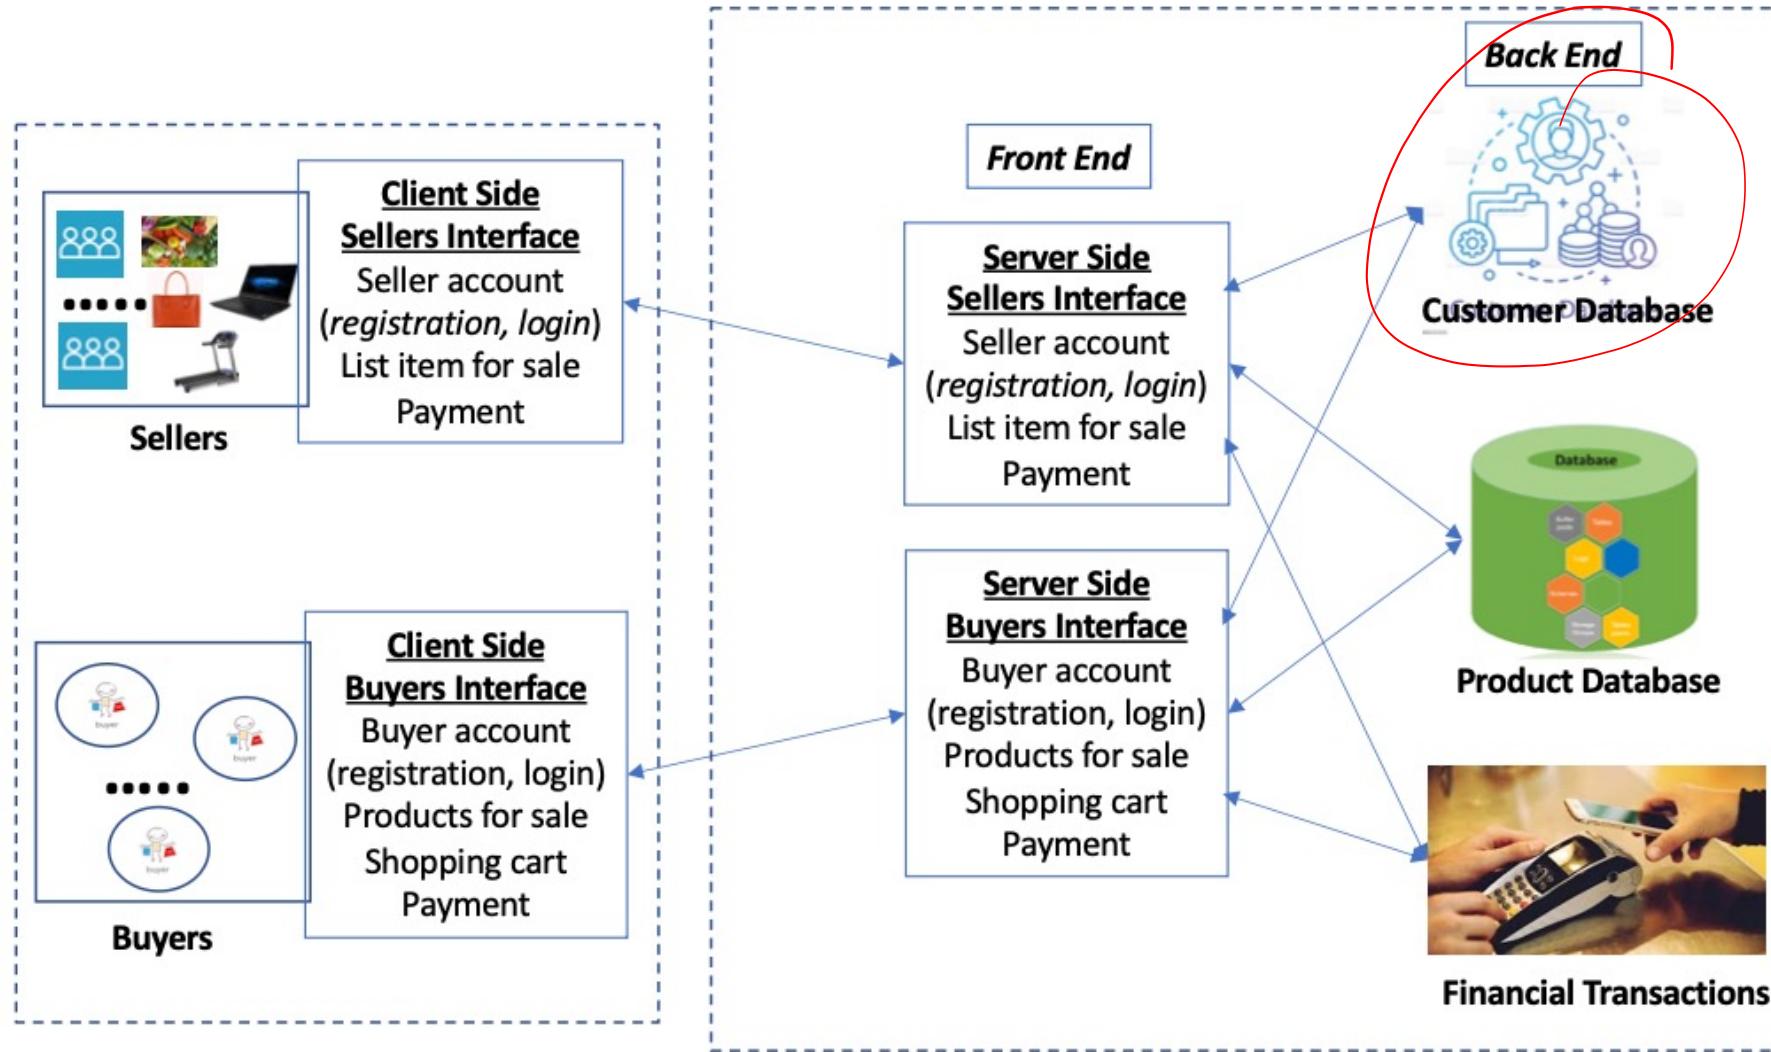
Electronic Marketplace



Interprocess Communication

- TCP/IP
 - Foundation of communication across distributed processes
 - Typically, too “low-level” to use directly
- Remote Procedure Calls (RPC)
 - Goal: Provide easy-to-program client/server communication
 - Review local procedure calls
 - Stub functions, parameter passing, server binding, data representation

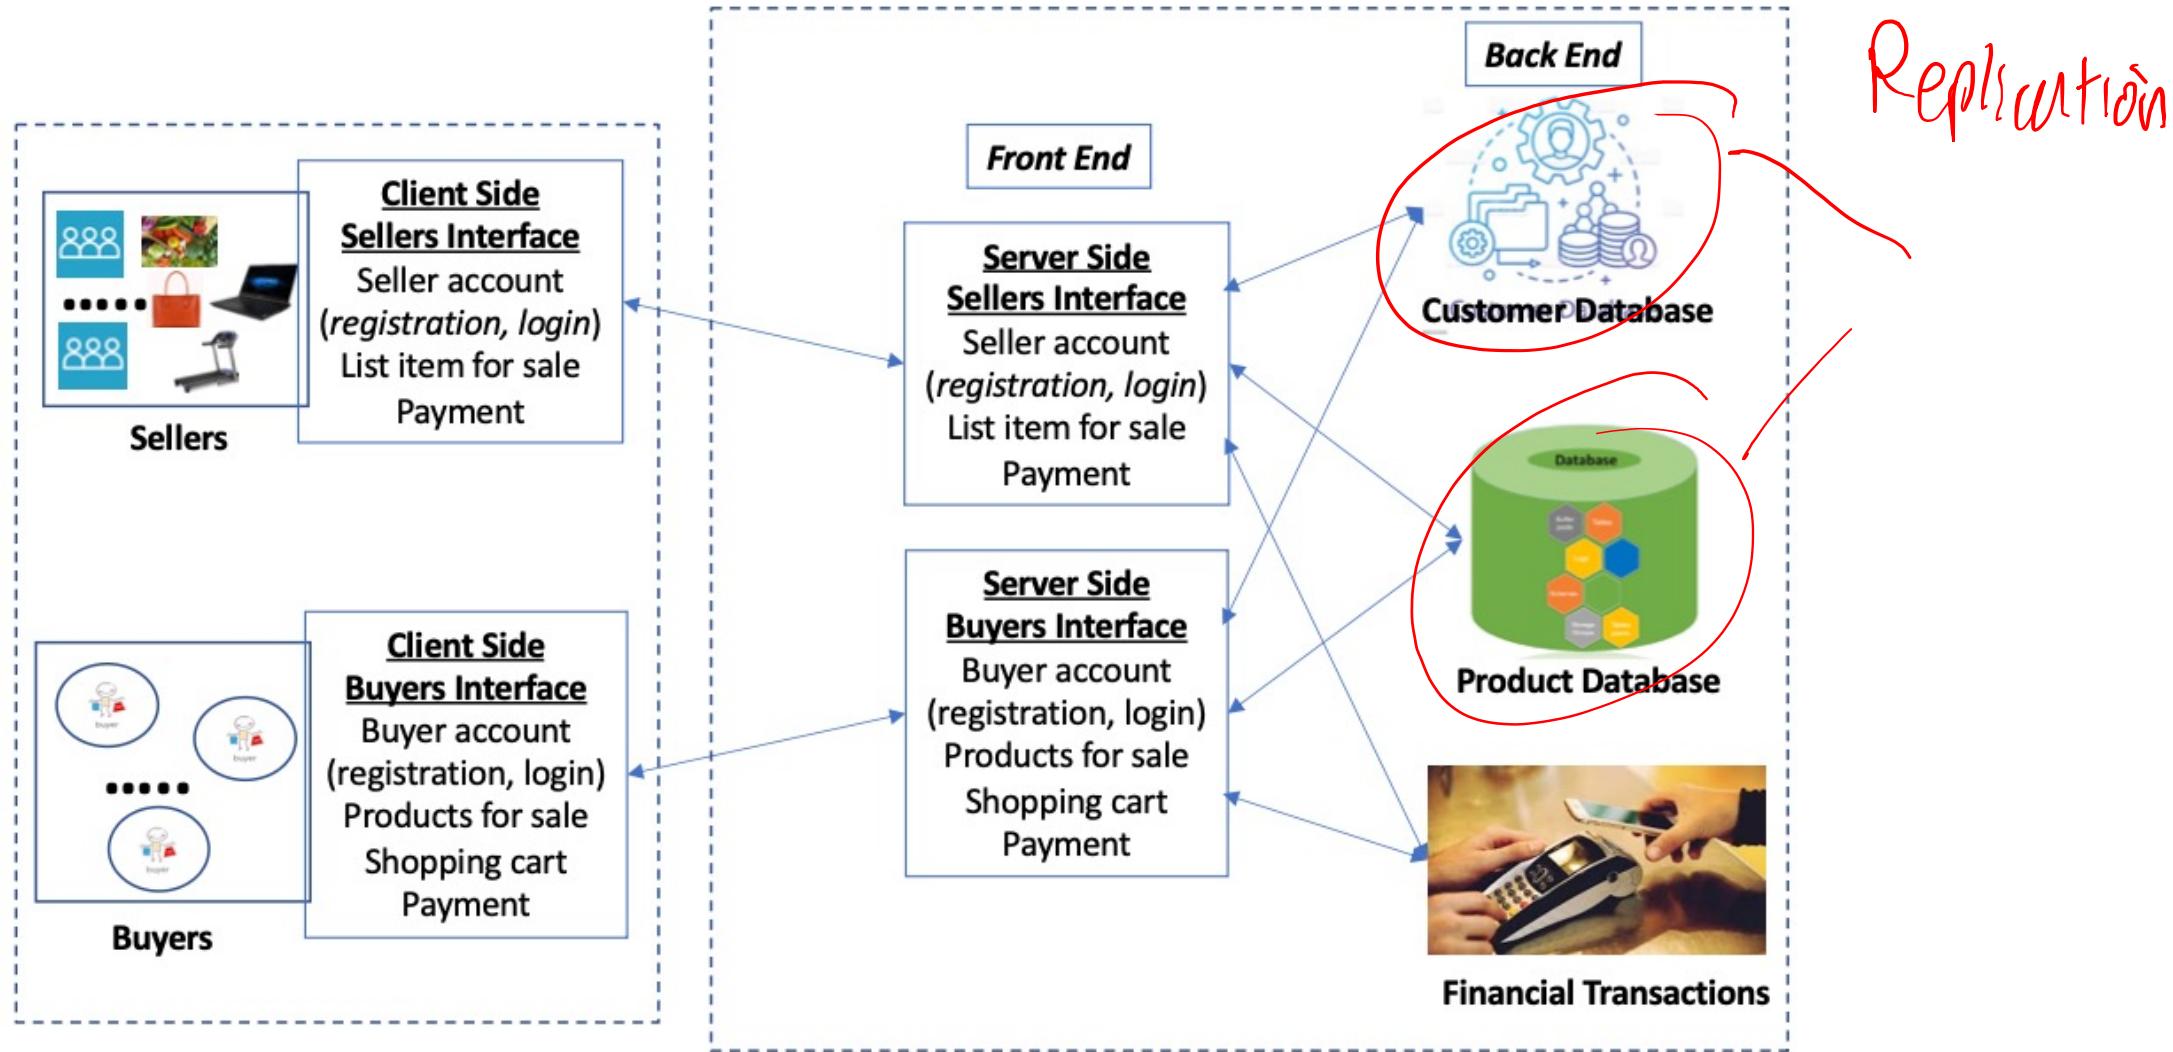
Electronic Marketplace



System models and failure models

- What do we care about in a distributed system?
 - What goals are we trying to achieve?
 - What assumptions can/can't we make?
 - What are some typical architectural styles?
- What can go wrong?
 - How can processes misbehave?
 - How can the network misbehave?
 - What can/can't we guarantee?

Electronic Marketplace



Managing replication

- Clock synchronization
 - How do we get clocks to agree as best as we can?
- Event ordering
 - How do we reason about the *relative* order of events across distributed processes?
 - Logical clocks
 - Vector clocks

Managing replication

- Communication and coordination
 - How do distributed replicas talk to each other?
 - Group communication systems
 - Leader election
- Consensus
 - What if things go wrong? How do we still make progress?
 - Byzantine Generals problem
 - Replicated state machines
 - Paxos, Raft
- Consistency and CAP
 - How do we reason about consistency across replicas?
 - CAP theorem – consistency, availability, and partition tolerance

Large-scale distributed systems

- How do we actually apply these principles into practice?
- How do we build a distributed locking and coordination service?
 - ZooKeeper, Chubby
- How do we process petabytes of data?
 - MapReduce, Spark
- How do we store petabytes of data?
 - Google File System, Meta's Tectonic File System

Large-scale distributed systems

- How do we build data abstractions on top of distributed storage/processing systems?
 - Amazon's Dynamo KV store
 - Cassandra database management system
 - Google's Spanner – globally distributed database
- Other topics
 - Peer-to-peer systems
 - “Serverless” computing
 - Machine learning systems